

Final Proj

May 8, 2024

How do specific game features and elements influence player satisfaction and engagement as reflected in review texts and ratings?

This question is aimed at identifying the key factors within games that contribute to or detract from player satisfaction. By analyzing the text of reviews in conjunction with the ratings, you can uncover which aspects of a game (like graphics, story, gameplay mechanics, or multiplayer features) are most closely associated with higher or lower customer ratings.

H1: Specific game features (such as story depth, graphical quality, or ease of controls) are positively correlated with high ratings.

H0: Game features do not have a significant impact on the ratings.

Feature Extraction: Identify mentions of specific game features in the review texts, categorizing them into aspects like story, graphics, gameplay, sound, and multiplayer functionality. **Sentiment Analysis:** Perform sentiment analysis on mentions of these features to determine whether the feedback is positive, negative, or neutral. **Correlation Analysis:** Analyze the correlation between the presence and sentiment of these features in reviews and the overall rating of the game.

Potential Business Impact: **Product Development:** Insights from this analysis can guide game developers in focusing their resources on enhancing the features that players care about the most. **Marketing Strategy:** Marketing teams can highlight the most praised features in their campaigns to attract a larger audience based on empirical data. **Customer Satisfaction:** By addressing the commonly criticized aspects, developers can improve player satisfaction and potentially increase positive word-of-mouth and ratings.

```
[9]: import nltk
from nltk import FreqDist
nltk.download("stopwords")
import pandas as pd
pd.set_option("display.max_colwidth",200)
import numpy as np
import re
import spacy
import matplotlib.pyplot as plt
import seaborn as sns
import json
import pandas as pd
import nltk
from nltk.corpus import stopwords
```

```

from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer # Import the stemmer
from tqdm import tqdm
%matplotlib inline

```

```

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\ghlas\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```

```

[3]: # Path to your JSON file
file_path = "S:\Python_Proj\Video_Games_5.json"

# Function to load and clean data
def load_and_clean_data(file_path):
    data = []
    with open(file_path, 'r', encoding='utf-8') as file:
        for line in file:
            try:
                # Load line as JSON and append to list
                entry = json.loads(line)
                # Remove unwanted columns
                entry.pop('unixReviewTime', None)
                entry.pop('vote', None)
                entry.pop('style', None)
                entry.pop('image', None)
                entry.pop('reviewTime', None)
                entry.pop('reviewerName', None)
                entry.pop('verified', None)
                entry.pop('summary', None)
                data.append(entry)
            except json.JSONDecodeError:
                continue
    return pd.DataFrame(data)

# Load the dataset and clean it
df = load_and_clean_data(file_path)
print(df.head())

```

```

<>:2: SyntaxWarning: invalid escape sequence '\P'
<>:2: SyntaxWarning: invalid escape sequence '\P'
C:\Users\ghlas\AppData\Local\Temp\ipykernel_31372\2746396954.py:2:
SyntaxWarning: invalid escape sequence '\P'
    file_path = "S:\Python_Proj\Video_Games_5.json"

```

	overall	reviewerID	asin \
0	5.0	A1HP7NVNPFMA4N	0700026657
1	4.0	A1JGAP0185YJI6	0700026657
2	3.0	A1YJWEXHQBWK2B	0700026657
3	2.0	A2204E1TH211HT	0700026657

4 5.0 A2RF5B5H74JLPE 0700026657

reviewText

0

This game is a bit hard to get the hang of, but when you do it's great.

1 I played it a while but it was alright. The steam was a bit of trouble. The more they move these game to steam the more of a hard time I have activating and playing a game. But in spite of that it...

2

ok game.

3

found the game a bit too complicated, not what I expected after having played 1602, 1503, and 1701

4

great game, I love it and have played it since its arrived

```
[4]: import pandas as pd

# Drop rows where 'reviewText' or 'overall' might be missing
df.dropna(subset=['reviewText', 'overall'], inplace=True)

# Explicitly convert review texts to string to avoid any confusion
df['reviewText'] = df['reviewText'].astype(str)
df['reviewText'] = df['reviewText'].str.replace(r"^[a-zA-Z\s]", "", regex=True).
    ↪str.lower()

# Create a new column to store the length of each review
df['review_length'] = df['reviewText'].apply(len)
all_strings = all(isinstance(x, str) for x in df['reviewText'])
print("All entries are strings:", all_strings)

# Display basic information about the DataFrame to ensure cleanliness
print(df.info())

# Display basic statistics of the review lengths
print(df['review_length'].describe())

print("Reviews longer than 2000 characters:", (df['review_length'] >= 1000).
    ↪sum())
df = df[df['review_length'] >= 1000]

# Print updated DataFrame info to confirm changes
print(df.info())
```

All entries are strings: True

<class 'pandas.core.frame.DataFrame'>

Index: 497419 entries, 0 to 497576

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	overall	497419 non-null	float64
1	reviewerID	497419 non-null	object
2	asin	497419 non-null	object
3	reviewText	497419 non-null	object
4	review_length	497419 non-null	int64

dtypes: float64(1), int64(1), object(3)
memory usage: 22.8+ MB

```
None
count      497419.00000
mean        646.58007
std         1221.11558
min          0.00000
25%          55.00000
50%         203.00000
75%         685.00000
max        31885.00000
Name: review_length, dtype: float64
Reviews longer than 2000 characters: 90182
<class 'pandas.core.frame.DataFrame'>
Index: 90182 entries, 13 to 497575
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	overall	90182 non-null	float64
1	reviewerID	90182 non-null	object
2	asin	90182 non-null	object
3	reviewText	90182 non-null	object
4	review_length	90182 non-null	int64

dtypes: float64(1), int64(1), object(3)
memory usage: 4.1+ MB
None

```
[5]: # Create a new column to store the length of each review
df['review_length'] = df['reviewText'].apply(len)
print(df['review_length'].describe())
```

```
count      90182.000000
mean        2512.675722
std         1921.096473
min         1000.000000
25%         1331.000000
50%         1871.000000
75%         2971.000000
max         31885.000000
Name: review_length, dtype: float64
```

```
[6]: nltk.download('punkt')
nltk.download('stopwords')

# Load spaCy's English language model
nlp = spacy.load('en_core_web_sm', disable=['parser', 'ner'])

# Initialize the Porter stemmer (from previous code)
from nltk.stem import PorterStemmer
stemmer = PorterStemmer()

# Function to preprocess text using both NLTK and spaCy for lemmatization
def preprocess_text(text):
    # Convert text to lowercase
    text = text.lower()
    # Tokenize text using NLTK
    tokens = word_tokenize(text)
    # Remove stop words using NLTK
    stop_words = set(stopwords.words('english'))
    filtered_tokens = [word for word in tokens if word not in stop_words and
↳ word.isalpha()]

    # Create a spaCy document for lemmatization
    doc = nlp(' '.join(filtered_tokens))
    lemmatized = [token.lemma_ for token in doc]

    return ' '.join(lemmatized)

# tqdm
tqdm.pandas()

df['processed_reviewText'] = df['reviewText'].progress_apply(preprocess_text)
```

```
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\ghlas\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\ghlas\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
100%|
| 90182/90182 [13:46<00:00, 109.17it/s]
```

```
[7]: # Display the first few processed texts
df = df.dropna(subset=['processed_reviewText'])
df = df[df['processed_reviewText'] != "nan"]
print(df[['reviewText', 'processed_reviewText']].head())
```

```
reviewText \
13 im not quite finished with the games dirt tour mode but i believe ive
```

experienced the bulk of what the game has to offer and im happy to say that the game is indeed awesome great cars great trac...

16 this is a must have for any gamer codemasters really hit a home run on this one i would like to see f be this good using the ego engine and the physics model in this game is right on the money t...

18 update june \ndeeply disappointed at the lack of a nonlinear openworld environment but its the restrictive qte gameplay and finalkill cutscenes that kill the game for me i loath enforced qte when ...

19 i will open with the pros\nreplayability its hitman what did you expect you can go back and replay missions over and over to try for a cleaner quicker kill or just fly right through it in a weeke...

29 dirt was like this im becoming more more suspicious of games like this that require a continuous open internet to play it just seems like bait set out for some young naive prey to fall for do...

processed_reviewText

13 I m quite finished game dirt tour mode believe I ve experience bulk game offer I m happy say game indeed awesome great car great track race mode excellent gameplay graphic highlight race snow vari...

16 must gamer codemaster really hit home run one would like see f good use ego engine physics model game right money game reason everyone buy good graphic card afford use nvidia overclocked order wan...

18 update june deeply disappointed lack nonlinear openworld environment restrictive qte gameplay finalkill cutscene kill game loath enforce qte mission go awry you re try escape hail bullet happen br...

19 open pro replayability hitman expect go back replay mission try clean quick kill fly right weekend like length game listen reviews xbox official magazine say lengthy game finish one weekend friday...

29 dirt like I m become suspicious game like require continuous open internet play seem like bait set young naive prey fall anyone honestly believe personal information collect particular reason do n...

```
[11]: import matplotlib.pyplot as plt

fig, axs = plt.subplots(1, 3, figsize=(18, 6))

# Plot 1: Distribution of Ratings
df['overall'].value_counts().sort_index().plot(kind='bar', ax=axs[0])
axs[0].set_title('Distribution of Ratings')
axs[0].set_xlabel('Rating')
axs[0].set_ylabel('Number of Reviews')

# Plot 2: Review Length by Rating
df['review_length'] = df['reviewText'].apply(len)
df.boxplot(column='review_length', by='overall', grid=False, showfliers=False,
           ax=axs[1])
axs[1].set_title('Review Length by Rating')
```

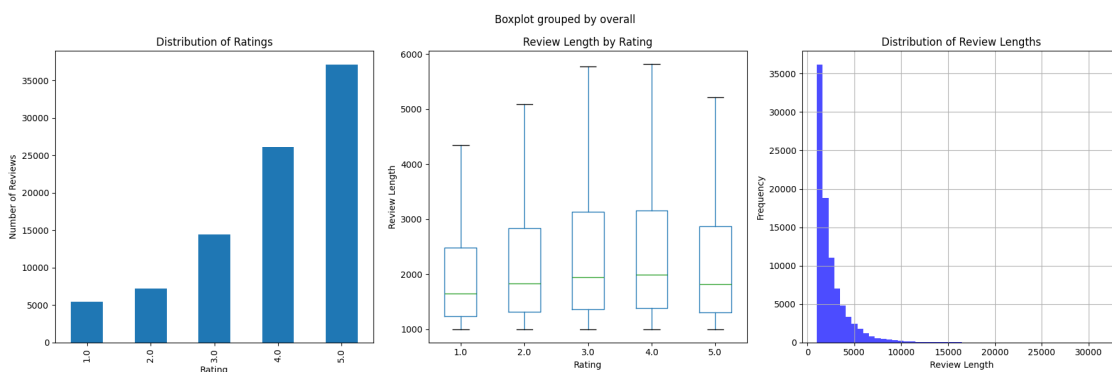
```

axs[1].set_xlabel('Rating')
axs[1].set_ylabel('Review Length')

# Plot 3: Distribution of Review Lengths
axs[2].hist(df['review_length'], bins=50, color='blue', alpha=0.7)
axs[2].set_title('Distribution of Review Lengths')
axs[2].set_xlabel('Review Length')
axs[2].set_ylabel('Frequency')
axs[2].grid(True)

fig.tight_layout()
plt.show()

```



```

[12]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import TruncatedSVD # Import TruncatedSVD for LSA
from tqdm.notebook import tqdm # tqdm tqdm_notebook

import pandas as pd

#
tfidf_vectorizer = TfidfVectorizer(max_features=4000, stop_words='english')

#
dtm = tfidf_vectorizer.fit_transform(df['processed_reviewText'])

#
n_topics = 5

# LSA
lsa = TruncatedSVD(n_components=n_topics, random_state=42)

# tqdm
tqdm.pandas()

```

```

# LSA
lsa.fit(dtm)

#
feature_names = tfidf_vectorizer.get_feature_names_out()
for topic_idx, topic in enumerate(tqdm(lsa.components_, desc='LSA Topic
↳Progress')):
    print(f"Topic {topic_idx}:")
    top_words_idx = topic.argsort()[::-1:-1]
    top_words = [feature_names[i] for i in top_words_idx]
    weights = [topic[i] for i in top_words_idx]
    print("\n".join([f"{word}: {weight}" for word, weight in zip(top_words,
↳weights)]))

```

LSA Topic Progress: 0% | 0/5 [00:00<?, ?it/s]

Topic 0:

game: 0.5236244122059027
play: 0.18016300135389915
like: 0.15390090028174647
time: 0.12125106375371393
good: 0.11774664023143717
character: 0.11651317305832855
make: 0.11146512667471252
really: 0.10702334038358467
story: 0.09953572496652699
fun: 0.09458125696402175

Topic 1:

controller: 0.37072029102441384
mouse: 0.33712381880502046
headset: 0.22997047448555222
button: 0.20074695751350788
use: 0.1702243690790324
ps: 0.16022573870845844
xbox: 0.15942834844076498
keyboard: 0.1370208459825826
work: 0.1167038592139729
wii: 0.11474022584685982

Topic 2:

mouse: 0.4852729581273917
enemy: 0.13780490998113784
character: 0.13449249645141031
keyboard: 0.12804342628630733
button: 0.11370904064892126
weapon: 0.1103817621713392
story: 0.10149601931635355
mission: 0.09549248470111553


```

key: 0.09542244075271651
battle: 0.08306624491810805
Topic 3:
mouse: 0.5003042692820453
mario: 0.29578819879605706
wii: 0.22715402697831813
game: 0.1505154629003359
nintendo: 0.12958982627848972
button: 0.11184146424872854
ds: 0.09865645348275648
super: 0.09307656643642526
kart: 0.0778941735869288
dpi: 0.07190828996205774
Topic 4:
car: 0.41408350559175955
race: 0.25864321152418523
sim: 0.18518189085723552
drive: 0.13558747056781076
mouse: 0.13182663469719985
mission: 0.11715130360674132
racing: 0.10110027212566493
online: 0.09427407254448114
mode: 0.08656292331903907
track: 0.08509581399701571

```

```

[13]: import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from tqdm.notebook import tqdm

# Define informative topics and their keywords
informative_topics = {
    'Graphics': [
        'resolution', 'textures', 'lighting', 'shadows', 'modeling',
        ↪ 'raytracing', 'HDR', 'framerate',
        'aliasing', 'rendering', '3D', 'detail', 'visuals', 'graphics',
        ↪ 'polygons', 'animation',
        'realism', 'artistic', 'effects', 'visual quality', 'color depth',
        ↪ 'bloom'
    ],
    'Story': [
        'plot', 'narrative', 'characters', 'dialogue', 'lore', 'cutscenes',
        ↪ 'script', 'storytelling',
        'development', 'twists', 'narration', 'themes', 'engagement', 'story',
        ↪ 'arcs', 'character arcs',

```

```

        'motivation', 'backstory', 'prologue', 'epilogue', 'drama', 'conflict',
        ↪ 'resolution'
    ],
    'Gameplay': [
        'mechanics', 'controls', 'design', 'difficulty', 'tutorials',
        ↪ 'interaction', 'AI', 'strategy',
        'levels', 'campaign', 'missions', 'challenges', 'playability', 'user_
        ↪ interface', 'HUD',
        'customization', 'skills', 'abilities', 'game balance', 'level_
        ↪ progression', 'sandbox',
        'open world', 'quest design'
    ],
    'Sound': [
        'soundtrack', 'effects', 'voices', 'audio', 'ambient', 'fidelity',
        ↪ 'acoustics', 'clarity',
        'mixing', '3Dsound', 'volume', 'music', 'sound design', 'dialogue_
        ↪ clarity', 'audio immersion',
        'surround sound', 'soundscapes', 'score', 'audio cues', 'background_
        ↪ noise', 'voiceovers'
    ],
    'Multiplayer': [
        'servers', 'matchmaking', 'co-op', 'pvp', 'esports', 'community',
        ↪ 'online', 'multiplayer',
        'teamplay', 'network', 'lobby', 'interaction', 'connectivity', 'social_
        ↪ features', 'clans',
        'guilds', 'teams', 'voice chat', 'competitive', 'leaderboards',
        ↪ 'rankings', 'cross-platform play'
    ]
}

```

```

[14]: text_data = df['processed_reviewText']
      # Vectorize the text data
      vectorizer = CountVectorizer(max_features=1000, stop_words='english')
      dtm = vectorizer.fit_transform(text_data)
      # Apply Latent Semantic Analysis (LSA)
      n_topics = 5
      lsa = TruncatedSVD(n_components=n_topics, random_state=42)
      lsa.fit(dtm)
      dtm_lsa = lsa.transform(dtm)
      # Normalize the LSA results
      dtm_lsa_normalized = normalize(dtm_lsa, axis=1)
      # Create a DataFrame to store LSA topic weights for each review
      df_lsa_topics = pd.DataFrame(dtm_lsa_normalized, columns=[f'LSA_Topic_{i}' for_
      ↪ i in range(n_topics)])

```

```
[15]: import pandas as pd
from tqdm.auto import tqdm
# DataFrame to store topic assignments based on keyword presence
df_informative_topics = pd.DataFrame(index=df.index)

# Iterate over each topic and check if any of its keywords are in the
↳processed_reviewText
for topic, keywords in tqdm(informative_topics.items(), desc="Processing
↳topics"):
    # Create a regex pattern that matches any of the keywords
    pattern = '|'.join([f'\\b{keyword}\\b' for keyword in keywords]) # \b is
↳used to ensure full words match
    df_informative_topics[topic] = df['processed_reviewText'].str.
↳contains(pattern, case=False, na=False)

print(df_informative_topics.head())
```

Processing topics: 0%| | 0/5 [00:00<?, ?it/s]

	Graphics	Story	Gameplay	Sound	Multiplayer
13	False	False	True	False	False
16	True	False	False	False	False
18	False	False	True	False	True
19	False	True	False	True	False
29	False	False	True	True	True

```
[16]: # Concatenate this DataFrame with the original DataFrame (or with LSA topics
↳DataFrame)
df_with_all_topics = pd.concat([df, df_lsa_topics, df_informative_topics],
↳axis=1)
df_with_all_topics = df_with_all_topics.dropna(subset=['processed_reviewText'])
df_with_all_topics =
↳df_with_all_topics[df_with_all_topics['processed_reviewText'] != "nan"]
print(df_with_all_topics[['reviewText', 'processed_reviewText']].head())
```

```

reviewText \
13 im not quite finished with the games dirt tour mode but i believe ive
experienced the bulk of what the game has to offer and im happy to say that the
game is indeed awesome great cars great trac...
16 this is a must have for any gamer codemasters really hit a home run on this
one i would like to see f be this good using the ego engine and the physics
model in this game is right on the money t...
18 update june \ndeeply disappointed at the lack of a nonlinear openworld
environment but its the restrictive qte gameplay and finalkill cutscenes that
kill the game for me i loath enforced qte when ...
19 i will open with the pros\nreplayability its hitman what did you expect you
can go back and replay missions over and over to try for a cleaner quicker kill
or just fly right through it in a weeke...
```

29 dirt was like this im becoming more more suspicious of games like this that require a continuous open internet to play it just seems like bait set out for some young naive prey to fall for do...

processed_reviewText

13 I m quite finished game dirt tour mode believe I ve experience bulk game offer I m happy say game indeed awesome great car great track race mode excellent gameplay graphic highlight race snow vari...

16 must gamer codemaster really hit home run one would like see f good use ego engine physics model game right money game reason everyone buy good graphic card afford use nvidia overclocked order wan...

18 update june deeply disappointed lack nonlinear openworld environment restrictive qte gameplay finalkill cutscene kill game loath enforce qte mission go awry you re try escape hail bullet happen br...

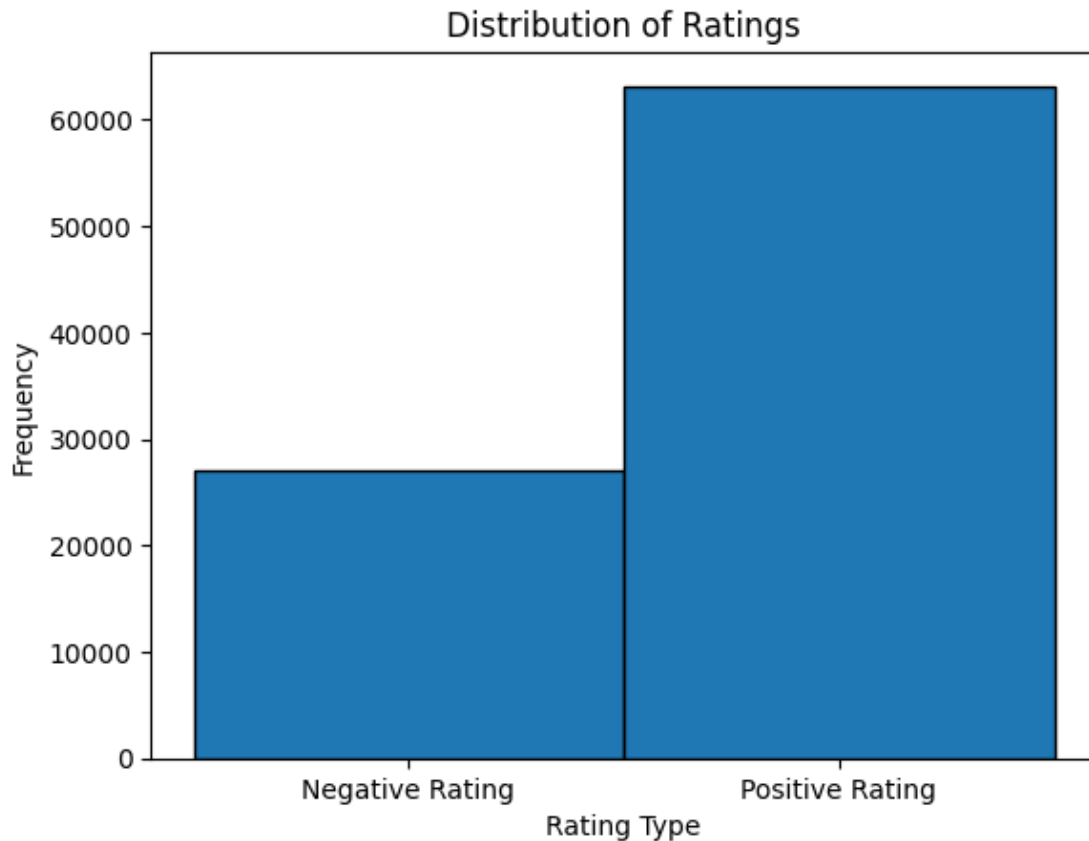
19 open pro replayability hitman expect go back replay mission try clean quick kill fly right weekend like length game listen reviews xbox official magazine say lengthy game finish one weekend friday...

29 dirt like I m become suspicious game like require continuous open internet play seem like bait set young naive prey fall anyone honestly believe personal information collect particular reason do n...

```
[17]: import numpy as np
import matplotlib.pyplot as plt

df_with_all_topics['overall_binary'] = np.where(df_with_all_topics['overall']_
    ↪ >= 4, 1, 0)

# Plotting the distribution of the new binary variable
plt.hist(df_with_all_topics['overall_binary'], bins=[-0.5, 0.5, 1.5],_
    ↪ edgecolor='black')
plt.xticks([0, 1], ['Negative Rating', 'Positive Rating'])
plt.xlabel('Rating Type')
plt.ylabel('Frequency')
plt.title('Distribution of Ratings')
plt.show()
```



```
[18]: from textblob import TextBlob

df_with_all_topics['processed_reviewText'] =
    ↪df_with_all_topics['processed_reviewText'].astype(str) # Ensure all data is
    ↪treated as string

# Function to compute sentiment from text
def compute_sentiment(text):
    if text.strip():
        return TextBlob(text).sentiment.polarity
    else:
        return 0

# Apply sentiment analysis
df_with_all_topics['sentiment'] = df_with_all_topics['processed_reviewText'].
    ↪apply(compute_sentiment)

# Display sentiment scores
print(df_with_all_topics[['processed_reviewText', 'sentiment']])
```

processed_reviewText \

13 I m quite finished game dirt tour mode believe I ve experience bulk game offer I m happy say game indeed awesome great car great track race mode excellent gameplay graphic highlight race snow vari...

16 must gamer codemaster really hit home run one would like see f good use ego engine physics model game right money game reason everyone buy good graphic card afford use nvidia overclocked order wan...

18 update june deeply disappointed lack nonlinear openworld environment restrictive qte gameplay finalkill cutscene kill game loath enforce qte mission go awry you re try escape hail bullet happen br...

19 open pro replayability hitman expect go back replay mission try clean quick kill fly right weekend like length game listen reviews xbox official magazine say lengthy game finish one weekend friday...

29 dirt like I m become suspicious game like require continuous open internet play seem like bait set young naive prey fall anyone honestly believe personal information collect particular reason do n...

...

...

497454 want thank reviewer talk game recommend long time think cheaply make first person parkour runner type game way wrong thinking wasent guy would never buy look ps exclusive play since everything ste...

497462 back original dishonor release dishonor everything fan would want sequel though small detail deep focus story make worthy successor franchise set several year event first game dishonor pick story ...

497538 buy game directly steam finish unlocked unlocked review pro definitely scare moment except first person view definitely old feel want hoard ammo stuff puzzle solve decent story decent visual excep...

497561 other may already state farm sim blast play busy busy sim farmer testify level authenticity say challenge nothing else always busy something think simfarm could reality play nonstop last four day ...

497575 think originally begin play bioshock several year ago be not level enough do not enough ammo remember reach point battle big daddy wrench fight ammo exhausted hour game do not want go back never p...


	sentiment
13	0.167435
16	0.066181
18	-0.046354
19	-0.038515
29	-0.009066
...	...
497454	0.095455
497462	0.109076
497538	0.046552
497561	0.081775
497575	0.015066

[90182 rows x 2 columns]

```
[19]: import seaborn as sns

# Calculate correlations between sentiment scores and overall ratings
correlations = df_with_all_topics[['sentiment', 'overall']].corr()

# Visualize the correlation matrix
sns.heatmap(correlations, annot=True, cmap='coolwarm')
plt.title('Correlation between Sentiment and Overall Ratings')
plt.show()
```



```
[32]: import seaborn as sns
import matplotlib.pyplot as plt

lsa_topics = ['LSA_Topic_0', 'LSA_Topic_1', 'LSA_Topic_2', 'LSA_Topic_3',
↳ 'LSA_Topic_4']
informative_topic_names = list(informative_topics.keys())

# Compute correlation matrices
lsa_correlation_matrix = df_with_all_topics[['overall', 'sentiment'] +
↳ lsa_topics].corr()
```

```

informative_correlation_matrix = df_with_all_topics[['overall', 'sentiment'] +
    ↪informative_topic_names].corr()


# Create subplots
fig, axs = plt.subplots(1, 2, figsize=(20, 8)) # 1 row, 2 columns

# Plot LSA Topic Correlation Matrix
sns.heatmap(lsa_correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
    ↪ax=axs[0])
axs[0].set_title('LSA Topic Correlation Matrix')

# Plot Informative Topic Correlation Matrix
sns.heatmap(informative_correlation_matrix, annot=True, cmap='coolwarm', fmt=".
    ↪2f", ax=axs[1])
axs[1].set_title('Informative Topic Correlation Matrix')

# Adjust layout for better spacing
fig.tight_layout()
plt.show()

```



```

[21]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from sklearn.metrics import accuracy_score, classification_report
import xgboost as xgb
from sklearn.metrics import accuracy_score

```



```

# Split the dataset
train, test = train_test_split(df_with_all_topics, test_size=0.2,
                                random_state=42)

# LSA topics
X_train_lsa = train[['LSA_Topic_0', 'LSA_Topic_1', 'LSA_Topic_2',
                    'LSA_Topic_3', 'LSA_Topic_4']]
y_train_lsa = train['overall_binary']
X_test_lsa = test[['LSA_Topic_0', 'LSA_Topic_1', 'LSA_Topic_2', 'LSA_Topic_3',
                  'LSA_Topic_4']]
y_test_lsa = test['overall_binary']

# Informative topics
informative_topic_names = list(informative_topics.keys())
X_train_info = train[informative_topic_names]
y_train_info = train['overall_binary']
X_test_info = test[informative_topic_names]
y_test_info = test['overall_binary']

def evaluate_model(model, X_train, y_train, X_test, y_test):
    model.fit(X_train, y_train)
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)
    accuracy_train = accuracy_score(y_train, y_pred_train)
    accuracy_test = accuracy_score(y_test, y_pred_test)
    print(classification_report(y_test, y_pred_test))
    return accuracy_train, accuracy_test

def evaluate_model_xgb(params, dtrain, dtest, num_rounds=100):
    # Train the model
    bst = xgb.train(params, dtrain, num_boost_round=num_rounds)

    # Make predictions
    y_pred_train = (bst.predict(dtrain) > 0.5).astype(int)
    y_pred_test = (bst.predict(dtest) > 0.5).astype(int)

    # Calculate accuracy
    accuracy_train = accuracy_score(dtrain.get_label(), y_pred_train)
    accuracy_test = accuracy_score(dtest.get_label(), y_pred_test)

    return accuracy_train, accuracy_test

```

```

[22]: import pandas as pd
import statsmodels.api as sm
import statsmodels.formula.api as smf
# Formula for LSA Topics

```

```

formula_lsa = 'overall_binary ~ LSA_Topic_0 + LSA_Topic_1 + LSA_Topic_2 +
↳LSA_Topic_3 + LSA_Topic_4'

# Fitting the OLS model
model_lsa = smf.ols(formula=formula_lsa, data=df_with_all_topics)
results_lsa = model_lsa.fit()
print("OLS Results for LSA Topics:")
print(results_lsa.summary())

# Formula for Informative Topics
formula_info = 'overall_binary ~ Graphics + Story + Gameplay + Sound +
↳Multiplayer'

# Fitting the OLS model
model_info = smf.ols(formula=formula_info, data=df_with_all_topics)
results_info = model_info.fit()
print("OLS Results for Informative Topics:")
print(results_info.summary())

# Logistic Regression for LSA Topics
model_lsa_logit = smf.logit(formula=formula_lsa, data=df_with_all_topics)
results_lsa_logit = model_lsa_logit.fit()
print("Logistic Regression Results for LSA Topics:")
print(results_lsa_logit.summary())

# Logistic Regression for Informative Topics
model_info_logit = smf.logit(formula=formula_info, data=df_with_all_topics)
results_info_logit = model_info_logit.fit()
print("Logistic Regression Results for Informative Topics:")
print(results_info_logit.summary())

```

OLS Results for LSA Topics:

OLS Regression Results						
=====						
Dep. Variable:	overall_binary	R-squared:	0.000			
Model:	OLS	Adj. R-squared:	-0.000			
Method:	Least Squares	F-statistic:	0.6253			
Date:	Wed, 08 May 2024	Prob (F-statistic):	0.680			
Time:	20:54:22	Log-Likelihood:	-15894.			
No. Observations:	27475	AIC:	3.180e+04			
Df Residuals:	27469	BIC:	3.185e+04			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	0.7974	0.032	24.839	0.000	0.735	0.860
LSA_Topic_0	-0.0491	0.035	-1.415	0.157	-0.117	0.019

LSA_Topic_1	-0.0189	0.015	-1.253	0.210	-0.049	0.011
LSA_Topic_2	0.0001	0.015	0.010	0.992	-0.030	0.030
LSA_Topic_3	0.0115	0.014	0.819	0.413	-0.016	0.039
LSA_Topic_4	-0.0142	0.017	-0.836	0.403	-0.047	0.019

Omnibus:	5123.426	Durbin-Watson:	1.779
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6725.031
Skew:	-1.170	Prob(JB):	0.00
Kurtosis:	2.370	Cond. No.	25.6

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Results for Informative Topics:

OLS Regression Results

Dep. Variable:	overall_binary	R-squared:	0.008
Model:	OLS	Adj. R-squared:	0.008
Method:	Least Squares	F-statistic:	147.8
Date:	Wed, 08 May 2024	Prob (F-statistic):	7.27e-157
Time:	20:54:22	Log-Likelihood:	-57171.
No. Observations:	90182	AIC:	1.144e+05
Df Residuals:	90176	BIC:	1.144e+05
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025
					0.975]

Intercept	0.6745	0.002	280.321	0.000	0.670
Graphics[T.True]	0.0403	0.004	10.661	0.000	0.033
Story[T.True]	-0.0201	0.004	-5.160	0.000	-0.028
Gameplay[T.True]	-0.0050	0.003	-1.565	0.118	-0.011
Sound[T.True]	0.0801	0.003	23.672	0.000	0.073
Multiplayer[T.True]	0.0002	0.003	0.052	0.959	-0.007

Omnibus:	102055.469	Durbin-Watson:	1.756
Prob(Omnibus):	0.000	Jarque-Bera (JB):	16774.887
Skew:	-0.864	Prob(JB):	0.00

Kurtosis: 1.784 Cond. No. 3.39

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Optimization terminated successfully.

Current function value: 0.559472

Iterations 5

Logistic Regression Results for LSA Topics:

Logit Regression Results

```
=====
Dep. Variable:    overall_binary    No. Observations:    27475
Model:            Logit             Df Residuals:        27469
Method:           MLE              Df Model:           5
Date:             Wed, 08 May 2024   Pseudo R-squ.:      0.0001020
Time:             20:54:22          Log-Likelihood:      -15371.
converged:        True              LL-Null:            -15373.
Covariance Type:  nonrobust         LLR p-value:         0.6790
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	1.3545	0.173	7.823	0.000	1.015	1.694
LSA_Topic_0	-0.2646	0.187	-1.416	0.157	-0.631	0.102
LSA_Topic_1	-0.1019	0.081	-1.254	0.210	-0.261	0.057
LSA_Topic_2	0.0008	0.082	0.009	0.993	-0.160	0.161
LSA_Topic_3	0.0622	0.076	0.821	0.412	-0.086	0.211
LSA_Topic_4	-0.0762	0.091	-0.837	0.402	-0.254	0.102

Optimization terminated successfully.

Current function value: 0.606178

Iterations 5

Logistic Regression Results for Informative Topics:

Logit Regression Results

```
=====
Dep. Variable:    overall_binary    No. Observations:    90182
Model:            Logit             Df Residuals:        90176
Method:           MLE              Df Model:           5
Date:             Wed, 08 May 2024   Pseudo R-squ.:      0.006836
Time:             20:54:23          Log-Likelihood:      -54666.
converged:        True              LL-Null:            -55043.
Covariance Type:  nonrobust         LLR p-value:         2.117e-160
=====
```

	coef	std err	z	P> z	[0.025	0.975]
--	------	---------	---	------	--------	--------

Intercept	0.7279	0.011	63.757	0.000	0.706
0.750					
Graphics[T.True]	0.1993	0.019	10.689	0.000	0.163
0.236					
Story[T.True]	-0.0969	0.019	-5.193	0.000	-0.134
-0.060					
Gameplay[T.True]	-0.0235	0.015	-1.546	0.122	-0.053
0.006					
Sound[T.True]	0.3984	0.017	23.555	0.000	0.365
0.432					
Multiplayer[T.True]	0.0010	0.017	0.059	0.953	-0.032
0.034					

=====

=====

```
[23]: from sklearn.model_selection import cross_val_score

# Random Forest - LSA
rf_lsa = RandomForestClassifier(n_estimators=50, random_state=42)
# Applying 5-fold Cross-Validation
cv_scores_lsa = cross_val_score(rf_lsa, X_train_lsa, y_train_lsa, cv=5,
                                ↪scoring='accuracy')

print("CV Scores for Random Forest LSA:", cv_scores_lsa)
print("Mean CV Accuracy for Random Forest LSA:", cv_scores_lsa.mean())

print("\n")

rf_info = RandomForestClassifier(n_estimators=50, random_state=42)
# Applying 5-fold Cross-Validation
cv_scores_info = cross_val_score(rf_info, X_train_info, y_train_info, cv=5,
                                  ↪scoring='accuracy')
# Print the CV scores
print("CV Scores for Random Forest Informative Topics:", cv_scores_info)
print("Mean CV Accuracy for Random Forest Informative Topics:", cv_scores_info.
      ↪mean())
```

CV Scores for Random Forest LSA: [0.6927715 0.69644466 0.69575161 0.69443482
0.69297942]

Mean CV Accuracy for Random Forest LSA: 0.694476401691039

CV Scores for Random Forest Informative Topics: [0.70108809 0.70108809
0.70101878 0.70101878 0.70101878]

Mean CV Accuracy for Random Forest Informative Topics: 0.701046503569201

```
[24]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay,
    ↪roc_curve, auc, precision_recall_curve, average_precision_score,
    ↪accuracy_score

# Function to perform CV and collect predictions
def cross_val_predict_proba(model, X, y, n_splits=5):
    kf = KFold(n_splits=n_splits, shuffle=True, random_state=42)
    y_trues, y_preds, y_scores = [], [], []

    for train_index, test_index in kf.split(X):
        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        model.fit(X_train, y_train)
        # Collect predictions
        y_pred = model.predict(X_test)
        y_prob = model.predict_proba(X_test)[:, 1] # assuming binary
    ↪classification

        y_trues.extend(y_test)
        y_preds.extend(y_pred)
        y_scores.extend(y_prob)

    return y_trues, y_preds, y_scores

# Plotting functions
def plot_evaluation_metrics(y_true, y_pred, y_scores, model_name):
    fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(18, 5))

    # Confusion Matrix
    ax1 = axes[0]
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(ax=ax1, cmap=plt.cm.Blues)
    ax1.set_title(f'Confusion Matrix for {model_name}')

    # ROC Curve
    ax2 = axes[1]
    fpr, tpr, _ = roc_curve(y_true, y_scores)
    roc_auc = auc(fpr, tpr)
    ax2.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area =
    ↪{roc_auc:.2f})')
    ax2.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
```

```

ax2.set_xlim([0.0, 1.0])
ax2.set_ylim([0.0, 1.05])
ax2.set_xlabel('False Positive Rate')
ax2.set_ylabel('True Positive Rate')
ax2.set_title(f'ROC Curve for {model_name}')
ax2.legend(loc="lower right")

# Precision-Recall Curve
ax3 = axes[2]
precision, recall, _ = precision_recall_curve(y_true, y_scores)
ap_score = average_precision_score(y_true, y_scores)
ax3.step(recall, precision, color='b', alpha=0.2, where='post')
ax3.fill_between(recall, precision, alpha=0.2, color='b')
ax3.set_xlabel('Recall')
ax3.set_ylabel('Precision')
ax3.set_ylim([0.0, 1.05])
ax3.set_xlim([0.0, 1.0])
ax3.set_title(f'Precision-Recall Curve for {model_name}: AP={ap_score:.2f}')

plt.tight_layout()
plt.show()

```

```

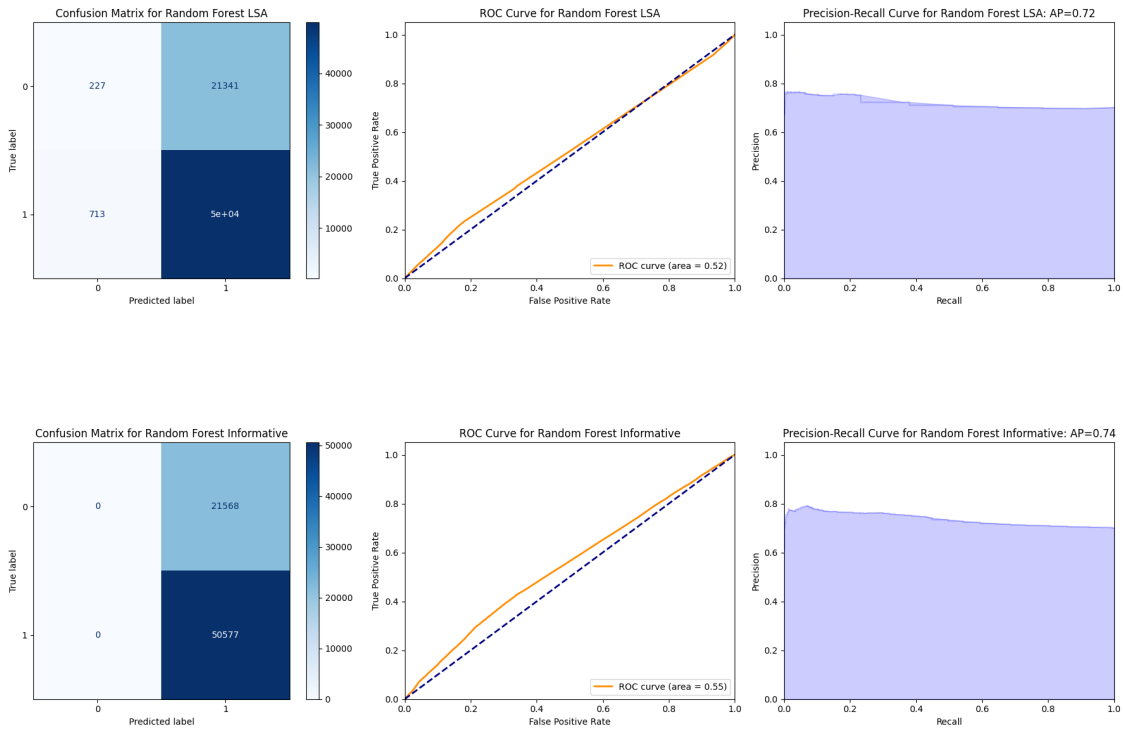
[25]: # Random Forest - LSA
rf_lsa = RandomForestClassifier(n_estimators=50, random_state=42)
y_true_lsa, y_pred_lsa, y_scores_lsa = cross_val_predict_proba(rf_lsa,
    ↪X_train_lsa, y_train_lsa)

# Random Forest - Informative Topics
rf_info = RandomForestClassifier(n_estimators=50, random_state=42)
y_true_info, y_pred_info, y_scores_info = cross_val_predict_proba(rf_info,
    ↪X_train_info, y_train_info)

# Plot metrics for LSA
plot_evaluation_metrics(y_true_lsa, y_pred_lsa, y_scores_lsa, "Random Forest_
    ↪LSA")

# Plot metrics for Informative Topics
plot_evaluation_metrics(y_true_info, y_pred_info, y_scores_info, "Random Forest_
    ↪Informative")

```



```
[26]: # Define a function to perform cross-validation and return average accuracy
def perform_cv(pipeline, X, y, cv=5, scoring='accuracy'):
    # Applying cross-validation
    cv_scores = cross_val_score(pipeline, X, y, cv=cv, scoring=scoring)

    # Return the mean of the cross-validation accuracy scores
    return cv_scores.mean()

# Neural Network - LSA Topics
pipeline_lsa_nn = make_pipeline(SimpleImputer(strategy='mean'),
    ↳MLPClassifier(hidden_layer_sizes=(100,), max_iter=500, random_state=42))
cv_accuracy_lsa_nn = perform_cv(pipeline_lsa_nn, X_train_lsa, y_train_lsa)
print("Neural Network LSA - CV Mean Accuracy:", cv_accuracy_lsa_nn)
print("\n")

# Neural Network - Informative Topics
pipeline_info_nn = make_pipeline(SimpleImputer(strategy='mean'),
    ↳MLPClassifier(hidden_layer_sizes=(100,), max_iter=500, random_state=42))
cv_accuracy_info_nn = perform_cv(pipeline_info_nn, X_train_info, y_train_info)
print("Neural Network Informative - CV Mean Accuracy:", cv_accuracy_info_nn)
```

Neural Network LSA - CV Mean Accuracy: 0.701046503569201

Neural Network Informative - CV Mean Accuracy: 0.701046503569201

```
[27]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, \
      ↪roc_curve, auc, precision_recall_curve, average_precision_score

def add_confusion_matrix(ax, y_true, y_pred, title):
    """Plot confusion matrix on a specific subplot axis."""
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(ax=ax, cmap=plt.cm.Blues)
    ax.title.set_text(title)

def add_roc_curve(ax, y_true, y_scores, title):
    """Plot ROC curve on a specific subplot axis."""
    fpr, tpr, _ = roc_curve(y_true, y_scores)
    roc_auc = auc(fpr, tpr)
    ax.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = \
    ↪{roc_auc:.2f})')
    ax.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    ax.set_xlim([0.0, 1.0])
    ax.set_ylim([0.0, 1.05])
    ax.set_xlabel('False Positive Rate')
    ax.set_ylabel('True Positive Rate')
    ax.legend(loc="lower right")
    ax.title.set_text(title)

def add_precision_recall_curve(ax, y_true, y_scores, title):
    """Plot precision-recall curve on a specific subplot axis."""
    precision, recall, _ = precision_recall_curve(y_true, y_scores)
    ap_score = average_precision_score(y_true, y_scores)
    ax.step(recall, precision, color='b', alpha=0.2, where='post')
    ax.fill_between(recall, precision, alpha=0.2, color='b')
    ax.set_xlabel('Recall')
    ax.set_ylabel('Precision')
    ax.set_ylim([0.0, 1.05])
    ax.set_xlim([0.0, 1.0])
    ax.title.set_text(f'Precision-Recall Curve: AP={ap_score:.2f} for {title}')

[28]: import matplotlib.pyplot as plt
from sklearn.pipeline import make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.neural_network import MLPClassifier

# Assuming models are trained; fit them if not already done
pipeline_lsa_nn = make_pipeline(SimpleImputer(strategy='mean'), \
    ↪MLPClassifier(hidden_layer_sizes=(100,), max_iter=500, random_state=42))
pipeline_lsa_nn.fit(X_train_lsa, y_train_lsa)
```

```

pipeline_info_nn = make_pipeline(SimpleImputer(strategy='mean'),
    ↳MLPClassifier(hidden_layer_sizes=(100,), max_iter=500, random_state=42))
pipeline_info_nn.fit(X_train_info, y_train_info)

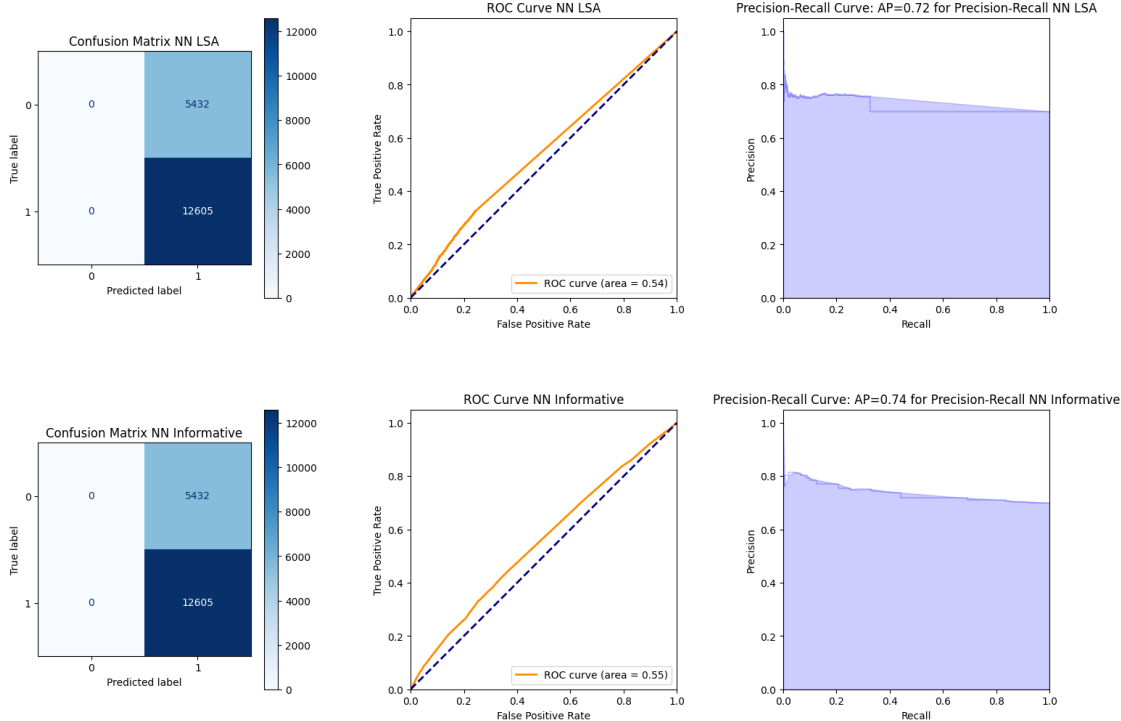
# Prepare the 2x3 grid of plots
fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(18, 12))
plt.subplots_adjust(hspace=0.4, wspace=0.4)

# Neural Network LSA Topics
y_pred_lsa_nn = pipeline_lsa_nn.predict(X_test_lsa)
y_scores_lsa_nn = pipeline_lsa_nn.predict_proba(X_test_lsa)[: , 1]
add_confusion_matrix(axes[0, 0], y_test_lsa, y_pred_lsa_nn, 'Confusion Matrix_
    ↳NN LSA')
add_roc_curve(axes[0, 1], y_test_lsa, y_scores_lsa_nn, 'ROC Curve NN LSA')
add_precision_recall_curve(axes[0, 2], y_test_lsa, y_scores_lsa_nn,
    ↳'Precision-Recall NN LSA')

# Neural Network Informative Topics
y_pred_info_nn = pipeline_info_nn.predict(X_test_info)
y_scores_info_nn = pipeline_info_nn.predict_proba(X_test_info)[: , 1]
add_confusion_matrix(axes[1, 0], y_test_info, y_pred_info_nn, 'Confusion Matrix_
    ↳NN Informative')
add_roc_curve(axes[1, 1], y_test_info, y_scores_info_nn, 'ROC Curve NN_
    ↳Informative')
add_precision_recall_curve(axes[1, 2], y_test_info, y_scores_info_nn,
    ↳'Precision-Recall NN Informative')

plt.show()

```



```
[29]: from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.svm import SVC

pipeline_lsa = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('svc', SVC(kernel='linear'))
])

pipeline_info = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')),
    ('svc', SVC(kernel='linear'))
])
```

```
[30]: # Fit the pipelines
pipeline_lsa.fit(X_train_lsa, y_train_lsa)
pipeline_info.fit(X_train_info, y_train_info)
```

```
[30]: Pipeline(steps=[('imputer', SimpleImputer()), ('svc', SVC(kernel='linear'))])
```

```
[34]: from sklearn.pipeline import make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
```

```

pipeline_lsa_svc = make_pipeline(
    SimpleImputer(strategy='mean'),
    StandardScaler(),
    SVC(kernel='linear', probability=True, random_state=42)
)
pipeline_info_svc = make_pipeline(
    SimpleImputer(strategy='mean'),
    StandardScaler(),
    SVC(kernel='linear', probability=True, random_state=42)
)

pipeline_lsa_svc.fit(X_train_lsa, y_train_lsa)
pipeline_info_svc.fit(X_train_info, y_train_info)


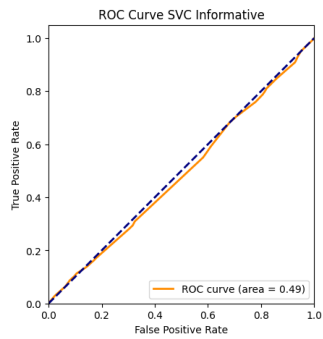


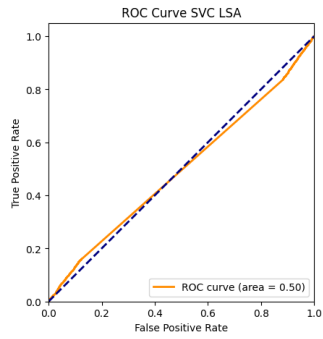

fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(18, 12))
plt.subplots_adjust(hspace=0.4, wspace=0.4)

# Predictions and evaluation for LSA SVC
y_pred_lsa_svc = pipeline_lsa_svc.predict(X_test_lsa)
y_scores_lsa_svc = pipeline_lsa_svc.predict_proba(X_test_lsa)[: , 1]
add_confusion_matrix(axes[0, 0], y_test_lsa, y_pred_lsa_svc, 'Confusion Matrix_
↳SVC LSA')
add_roc_curve(axes[0, 1], y_test_lsa, y_scores_lsa_svc, 'ROC Curve SVC LSA')
add_precision_recall_curve(axes[0, 2], y_test_lsa, y_scores_lsa_svc,
↳'Precision-Recall SVC LSA')

# Predictions and evaluation for Informative SVC
y_pred_info_svc = pipeline_info_svc.predict(X_test_info)
y_scores_info_svc = pipeline_info_svc.predict_proba(X_test_info)[: , 1]
add_confusion_matrix(axes[1, 0], y_test_info, y_pred_info_svc, 'Confusion_
↳Matrix SVC Informative')
add_roc_curve(axes[1, 1], y_test_info, y_scores_info_svc, 'ROC Curve SVC_
↳Informative')
add_precision_recall_curve(axes[1, 2], y_test_info, y_scores_info_svc,
↳'Precision-Recall SVC Informative')

plt.show()

```



[]:

[]: